


1. Primary paradox test (the central causal claim)



Paradox: higher predictive fidelity (calibrated GB, matched-resolution BB) collapses in live control ($m_s > 1$); only the coarse hourly BB and the role-separated hybrid stay usable.

2. Boundary checks (scope of the recipe — not part of the central test)

HDRL

censor weight is controller-family specific ($\lambda_T=0$)

MORL

viable only with 17D forecast observation (5D fails)

Hydronic transfer (3 cases)

calibration transfers;
frozen policy does not

Live metric: m_s = time-in-violation + worst-case severity ($m_s > 1$ = unusable); also comfort-violation %, energy, and 24 h rollout RMSE_T .