


MORL: preference-conditioned four-stage pipeline

17D observation + power-only hybrid backend ($\lambda_T = 0$, $\lambda_P = 5 \times 10^{-5}$)



preference weights $w = (w_c, w_e, w_s)$ condition the scalar reward $r = w_c r^c + w_e r^e + w_s r^s$;

MORL is the only family with a live-BOPTEST finetune; PPO/HDRL are evaluated zero-shot.